

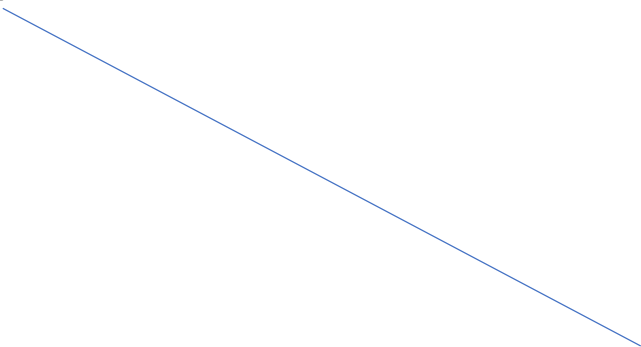
Gradient Descent

Dr. Xiaowei Huang

<https://cgi.csc.liv.ac.uk/~xiaowei/>

Up to now,

- Three machine learning algorithms:
 - decision tree learning
 - k-nn
 - linear regression



only optimization
objectives are
discussed, but
how to solve?

Today's Topics

- Derivative
- Gradient
- Directional Derivative
- Method of Gradient Descent
- Example: Gradient Descent on Linear Regression
- Linear Regression: Analytical Solution

Problem Statement: Gradient-Based Optimization

- Most ML algorithms involve optimization
- Minimize/maximize a function $f(\mathbf{x})$ by altering \mathbf{x}
 - Usually stated as a minimization of e.g., the loss etc
 - Maximization accomplished by minimizing $-f(\mathbf{x})$
- $f(\mathbf{x})$ referred to as objective function or criterion
 - In minimization also referred to as loss function cost, or error
 - Example:
 - linear least squares $f(x) = \frac{1}{2} \|Ax - b\|^2$
 - **Linear regression** $\hat{L}(f_w) = \frac{1}{m} \sum_{i=1}^m (w^T x^{(i)} - y^{(i)})^2$
 - Denote optimum value by $\mathbf{x}^* = \operatorname{argmin} f(\mathbf{x})$


Derivative

Derivative of a function

- Suppose we have function $y=f(x)$, x, y real numbers
 - Derivative of function denoted: $f'(x)$ or as dy/dx
 - Derivative $f'(x)$ gives the slope of $f(x)$ at point x
 - It specifies how to scale a small change in input to obtain a corresponding change in the output:

$$f(x + \varepsilon) \approx f(x) + \varepsilon f'(x)$$

- It tells how you make a small change in input to make a small improvement in y



Recall what's the derivative for the following functions:

$$f(x) = x^2$$

$$f(x) = e^x$$

...

Calculus in Optimization

- Suppose we have function $y = f(x)$, where x, y are real numbers

- Sign function:

$$\text{sign}(x) = \begin{cases} -1 & \text{if } x < 0 \\ 0 & \text{if } x = 0 \\ 1 & \text{if } x > 0 \end{cases}$$

- We know that

$$f(x - \epsilon \text{sign}(f'(x))) < f(x)$$

for small ϵ .


This technique is called **gradient descent** (Cauchy 1847)

- Therefore, we can reduce $f(x)$ by moving x in small steps with **opposite** sign of derivative


Why opposite?

Example

- Function $f(x) = x^2$ $\varepsilon = 0.1$
- $f'(x) = 2x$
- For $x = -2$, $f'(-2) = -4$, $\text{sign}(f'(-2)) = -1$
- $f(-2 - \varepsilon * (-1)) = f(-1.9) < f(-2)$
- For $x = 2$, $f'(2) = 4$, $\text{sign}(f'(2)) = 1$
- $f(2 - \varepsilon * 1) = f(1.9) < f(2)$



Gradient Descent Illustrated



For $x < 0$, $f(x)$ decreases with x
and $f'(x) < 0$


For $x > 0$, $f(x)$ increases with x
and $f'(x) > 0$

Use $f'(x)$ to follow
function downhill

Reduce $f(x)$ by going in direction
opposite sign of derivative $f'(x)$

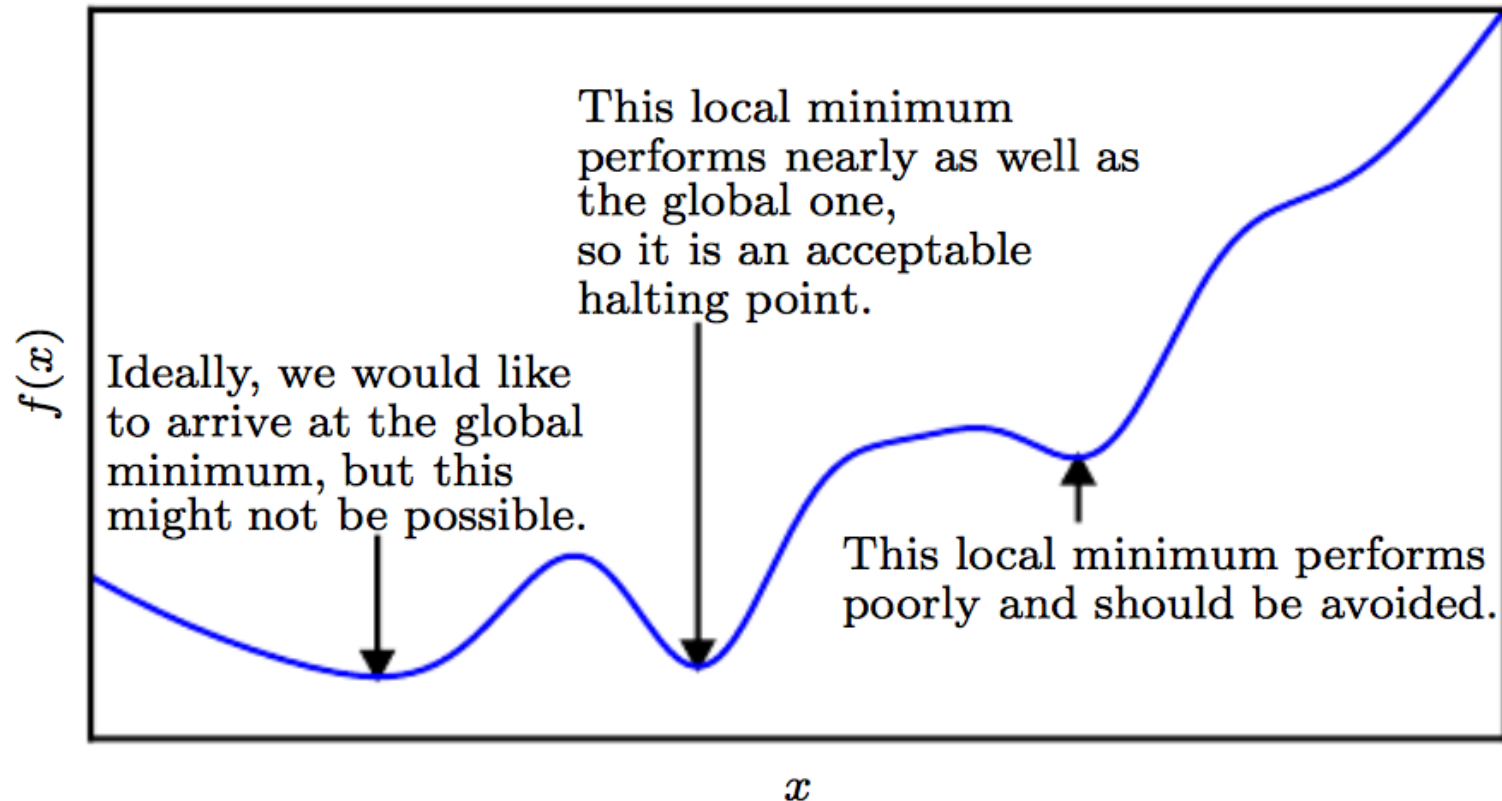
Stationary points, Local Optima

- When $f'(x) = 0$ derivative provides no information about direction of move
- Points where $f'(x) = 0$ are known as *stationary* or critical points
 - Local minimum/maximum: a point where $f(x)$ lower/ higher than all its neighbors
 - Saddle Points: neither maxima nor minima



Presence of multiple minima

- Optimization algorithms may fail to find global minimum
- Generally accept such solutions



Gradient

Minimizing with multiple dimensional inputs

- We often minimize functions with multiple-dimensional inputs

$$f : \mathbb{R}^n \rightarrow \mathbb{R}$$

- For minimization to make sense there must still be only one (scalar) output

Functions with multiple inputs

- Partial derivatives

$$\frac{\partial}{\partial x_i} f(x)$$

measures how f changes as only variable x_i increases at point \mathbf{x}

- Gradient generalizes notion of derivative where derivative is wrt a vector
- Gradient is vector containing all of the partial derivatives denoted

$$\nabla_x f(x) = \left(\frac{\partial}{\partial x_1} f(x), \dots, \frac{\partial}{\partial x_n} f(x) \right)$$

Example

- $y = 5x_1^5 + 4x_2 + x_3^2 + 2$
- so what is the exact gradient on instance (1,2,3)
- the gradient is $(25x_1^4, 4, 2x_3)$
- On the instance (1,2,3), it is (25,4,6)

Functions with multiple inputs

- Gradient is vector containing all of the partial derivatives denoted

$$\nabla_x f(x) = \left(\frac{\partial}{\partial x_1} f(x), \dots, \frac{\partial}{\partial x_n} f(x) \right)$$

- Element i of the gradient is the partial derivative of f wrt x_i
- Critical points are where every element of the gradient is equal to zero

$$\nabla_x f(x) = 0 \equiv \begin{cases} \frac{\partial}{\partial x_1} f(x) = 0 \\ \dots \\ \frac{\partial}{\partial x_n} f(x) = 0 \end{cases}$$

Example

- $y = 5x_1^5 + 4x_2 + x_3^2 + 2$
- so what are the critical points?
- the gradient is $(25x_1^4, 4, 2x_3)$
- We let $25x_1^4 = 0$ and $2x_3 = 0$, so all instances whose x_1 and x_3 are 0. but $4 \neq 0$. So there is no critical point.

Directional Derivative

Directional Derivative

- Directional derivative in direction \mathbf{u} (a unit vector) is the slope of function f in direction \mathbf{u}

- This evaluates to

$$\mathbf{u}^T \nabla_x f(x)$$

- Example: let $\mathbf{u}^T = (u_x, u_y, u_z)$ be a unit vector in Cartesian coordinates, so

$$\|\mathbf{u}\|_2 = \sqrt{u_x^2 + u_y^2 + u_z^2} = 1$$

then

$$\mathbf{u}^T \nabla_x f(x) = \frac{\partial f}{\partial x} u_x + \frac{\partial f}{\partial y} u_y + \frac{\partial f}{\partial z} u_z$$

Directional Derivative

- To minimize f find direction in which f decreases the fastest

$$\min_{u, u^T u = 1} u^T \nabla_x f(x) = \min_{u, u^T u = 1} \|u\|_2 \cdot \|\nabla_x f(x)\|_2 \cdot \cos \theta$$

- where θ is angle between u and the gradient
- Substitute $\|u\|_2 = 1$ and ignore factors that not depend on u this simplifies to

$$\min_u \cos \theta$$

- This is minimized when u points in direction opposite to gradient
- In other words, the *gradient points directly uphill, and the negative gradient points directly downhill*

Method of Gradient Descent

Method of Gradient Descent

- The gradient points directly uphill, and the negative gradient points directly downhill
- Thus we can decrease f by moving in the direction of the negative gradient
 - This is known as the method of **steepest descent or gradient descent**
- Steepest descent proposes a new point

$$x' = x - \epsilon \nabla_x f(x)$$

- where ϵ is the learning rate, a positive scalar. Set to a small constant.

Choosing ϵ : Line Search

- We can choose ϵ in several different ways
- Popular approach: set ϵ to a small constant
- Another approach is called *line search*:
 - Evaluate

$$f(x - \epsilon \nabla_x f(x))$$

for several values of ϵ and choose the one that results in smallest objective function value

Example: Gradient Descent on Linear Regression

Example: Gradient Descent on Linear Regression

- Linear regression: $\hat{L}(f_w) = \frac{1}{m} \sum_{i=1}^m (w^T x^{(i)} - y^{(i)})^2 = \frac{1}{m} \|Xw - y\|_2^2$
- The gradient is

$$\begin{aligned} & \nabla_w \hat{L}(f_w) \\ = & \nabla_w \frac{1}{m} \|Xw - y\|_2^2 \\ = & \nabla_w [(Xw - y)^T (Xw - y)] \\ = & \nabla_w [w^T X^T Xw - 2w^T X^T y + y^T y] \\ = & 2X^T Xw - 2X^T y \end{aligned}$$

Example: Gradient Descent on Linear Regression

- Linear regression: $\hat{L}(f_w) = \frac{1}{m} \sum_{i=1}^m (w^T x^{(i)} - y^{(i)})^2 = \frac{1}{m} \|Xw - y\|_2^2$
- The gradient is $\nabla_w \hat{L}(f_w) = 2X^T Xw - 2X^T y$
- Gradient Descent algorithm is
 - Set step size ϵ , tolerance δ to small, positive numbers.
 - *While* $\|X^T Xw - X^T y\|_2 > \delta$ *do*

$$x \longleftarrow x - \epsilon(X^T Xw - X^T y)$$

Linear Regression: Analytical solution

Convergence of Steepest Descent

- Steepest descent converges when every element of the gradient is zero
 - In practice, very close to zero
- We may be able to avoid iterative algorithm and jump to the critical point by solving the following equation for x

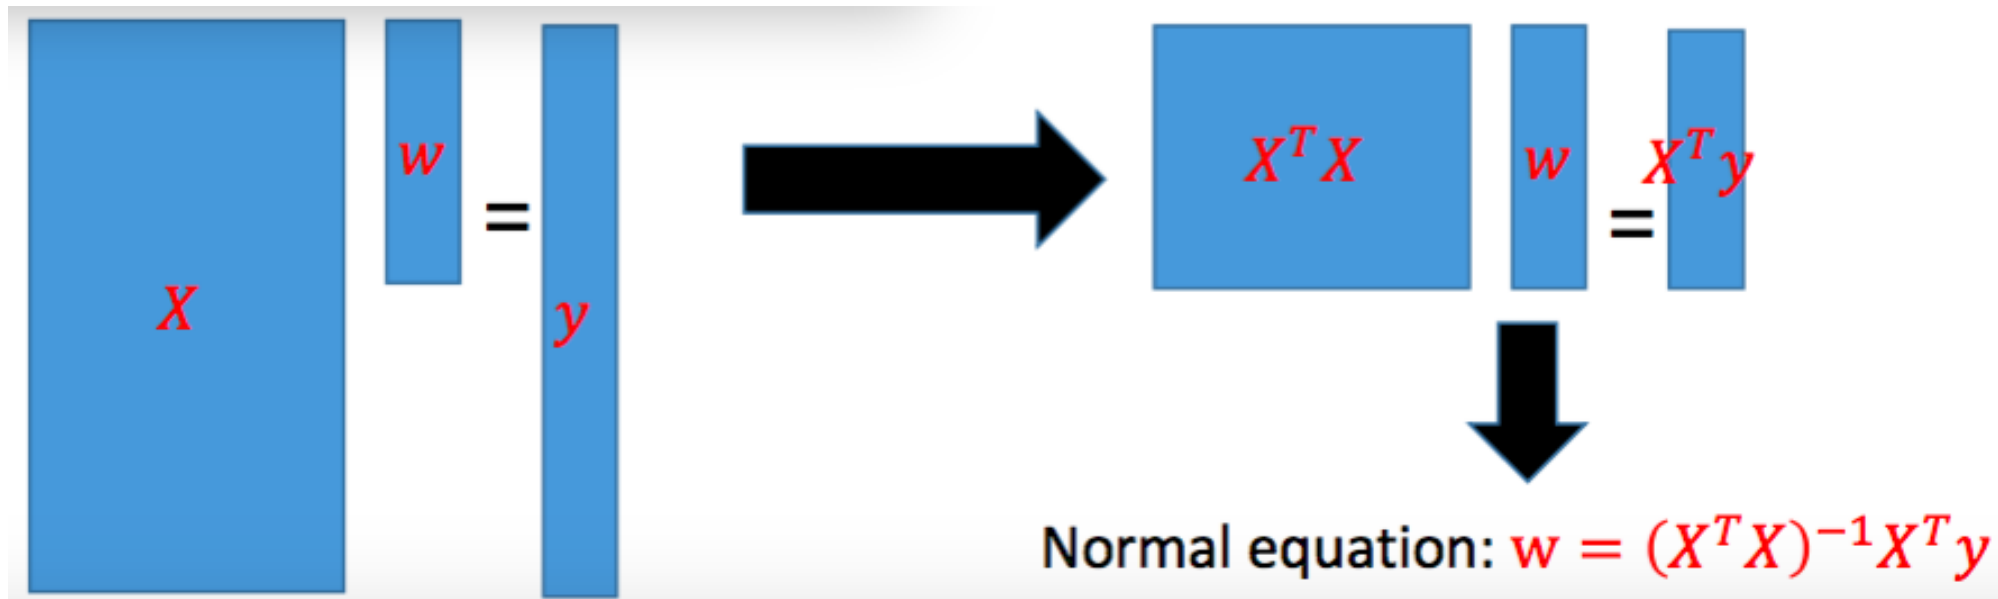
$$\nabla_x f(x) = 0$$

Linear Regression: Analytical solution

- Linear regression: $\hat{L}(f_w) = \frac{1}{m} \sum_{i=1}^m (w^T x^{(i)} - y^{(i)})^2 = \frac{1}{m} \|Xw - y\|_2^2$
- The gradient is $\nabla_w \hat{L}(f_w) = 2X^T Xw - 2X^T y$
- Let $\nabla_w \hat{L}(f_w) = 2X^T Xw - 2X^T y = 0$
- Then, we have $w = (X^T X)^{-1} X^T y$

Linear Regression: Analytical solution

- Algebraic view of the minimizer
- If X is invertible, just solve $Xw = y$ and get $w = X^{-1}y$
- But typically X is a tall matrix



Generalization to discrete spaces

Generalization to discrete spaces

- Gradient descent is limited to continuous spaces
- Concept of repeatedly making the best small move can be generalized to discrete spaces
- Ascending an objective function of discrete parameters is called *hill climbing*

Exercises

- Given a function $f(x) = e^x / (1 + e^x)$, how many critical points?
- Given a function $f(x_1, x_2) = 9x_1^2 + 3x_2 + 4$, how many critical points?
- Please write a program to do the following: given any differentiable function (such as the above two), an ε , and a starting x and a target x' , determine whether it is possible to reach x' from x . If possible, how many steps? You can adjust ε to see the change of the answer.

Extended Materials

Beyond Gradient: Jacobian and Hessian matrices

- Sometimes we need to find all derivatives of a function whose input and output are both vectors
- If we have function $f: R_m \rightarrow R_n$
 - Then the matrix of partial derivatives is known as the Jacobian matrix J defined as


$$J_{i,j} = \frac{\partial}{\partial x_j} f(x)_i$$

Second derivative

- Derivative of a derivative
- For a function $f: R^n \rightarrow R$ the derivative wrt x_i of the derivative of f wrt x_j is denoted as $\frac{\partial^2}{\partial x_i \partial x_j} f$
- In a single dimension we can denote $\frac{\partial^2}{\partial x^2} f$ by $f''(\mathbf{x})$
- Tells us how the first derivative will change as we vary the input
- This is important as it tells us whether a gradient step will cause as much of an improvement as based on gradient alone

Second derivative measures curvature

- Derivative of a derivative
- Quadratic functions with different curvatures



Hessian

- Second derivative with many dimensions

- $H(f)(x)$ is defined as

$$H(f)(x)_{i,j} = \frac{\partial^2}{\partial x_i \partial x_j} f(x)$$

- Hessian is the Jacobian of the gradient
- Hessian matrix is symmetric, i.e., $H_{i,j} = H_{j,i}$
 - anywhere that the second partial derivatives are continuous
 - So the Hessian matrix can be decomposed into a set of real eigenvalues and an orthogonal basis of eigenvectors
 - Eigenvalues of H are useful to determine learning rate as seen in next two slides

Role of eigenvalues of Hessian

- Second derivative in direction d is $d^T H d$
 - If d is an eigenvector, second derivative in that direction is given by its eigenvalue
 - For other directions, weighted average of eigenvalues (weights of 0 to 1, with eigenvectors with smallest angle with d receiving more value)
- Maximum eigenvalue determines maximum second derivative and minimum eigenvalue determines minimum second derivative

Learning rate from Hessian

- Taylor's series of $f(\mathbf{x})$ around current point $\mathbf{x}^{(0)}$

$$f(\mathbf{x}) \approx f(\mathbf{x}^{(0)}) + (\mathbf{x} - \mathbf{x}^{(0)})^T \mathbf{g} + \frac{1}{2}(\mathbf{x} - \mathbf{x}^{(0)})^T H(\mathbf{x} - \mathbf{x}^{(0)})$$

- where \mathbf{g} is the gradient and H is the Hessian at $\mathbf{x}^{(0)}$
- If we use learning rate ϵ the new point \mathbf{x} is given by $\mathbf{x}^{(0)} - \epsilon \mathbf{g}$. Thus we get

$$f(\mathbf{x}^{(0)} - \epsilon \mathbf{g}) \approx f(\mathbf{x}^{(0)}) - \epsilon \mathbf{g}^T \mathbf{g} + \frac{1}{2} \epsilon^2 \mathbf{g}^T H \mathbf{g}$$

- There are three terms:
 - original value of f ,
 - expected improvement due to slope, and
 - correction to be applied due to curvature
- Solving for step size when correction is least gives

$$\epsilon^* \approx \frac{\mathbf{g}^T \mathbf{g}}{\mathbf{g}^T H \mathbf{g}}$$

Second Derivative Test: Critical Points


- On a critical point $f'(x)=0$
- When $f''(x)>0$ the first derivative $f'(x)$ increases as we move to the right and decreases as we move left
- We conclude that x is a local minimum
- For local maximum, $f'(x)=0$ and $f''(x)<0$
- When $f''(x)=0$ test is inconclusive: x may be a saddle point or part of a flat region

Multidimensional Second derivative test

- In multiple dimensions, we need to examine second derivatives of all dimensions
- Eigendecomposition generalizes the test
- Test eigenvalues of Hessian to determine whether critical point is a local maximum, local minimum or saddle point
- When H is positive definite (all eigenvalues are positive) the point is a local minimum
- Similarly negative definite implies a maximum

Saddle point

- Contains both positive and negative curvature
- Function is $f(\mathbf{x}) = x_1^2 - x_2^2$



- Along axis x_1 , function curves upwards: this axis is an eigenvector of H and has a positive value
- Along x_2 , function curves downwards; its direction is an eigenvector of H with negative eigenvalue
- At a saddle point eigen values are both positive and negative

Inconclusive Second Derivative Test


- Multidimensional second derivative test can be inconclusive just like univariate case
- Test is inconclusive when all non-zero eigen values have same sign but at least one value is zero
 - since univariate second derivative test is inconclusive in cross-section corresponding to zero eigenvalue

Poor Condition Number

- There are different second derivatives in each direction at a single point
- Condition number of H e.g., $\lambda_{max}/\lambda_{min}$ measures how much they differ
 - Gradient descent performs poorly when H has a poor condition no.
 - Because in one direction derivative increases rapidly while in another direction it increases slowly
 - Step size must be small so as to avoid overshooting the minimum, but it will be too small to make progress in other directions with less curvature

Gradient Descent without H

- H with condition no, 5
 - Direction of most curvature has five times more curvature than direction of least curvature
- Due to small step size Gradient descent wastes time
- Algorithm based on Hessian can predict that steepest descent is not promising



Newton's method uses Hessian

- Another second derivative method
 - Using Taylor's series of $f(\mathbf{x})$ around current $\mathbf{x}^{(0)}$

$$f(\mathbf{x}) \approx f(\mathbf{x}^{(0)}) + (\mathbf{x} - \mathbf{x}^{(0)})^T \nabla_{\mathbf{x}} f(\mathbf{x}^{(0)}) + \frac{1}{2} (\mathbf{x} - \mathbf{x}^{(0)})^T H(f)(\mathbf{x} - \mathbf{x}^{(0)}) (\mathbf{x} - \mathbf{x}^{(0)})$$

- solve for the critical point of this function to give $\mathbf{x}^* = \mathbf{x}^{(0)} - H(f)(\mathbf{x}^{(0)})^{-1} \nabla_{\mathbf{x}} f(\mathbf{x}^{(0)})$
 - When f is a quadratic (positive definite) function use solution to jump to the minimum function directly
 - When not quadratic apply solution iteratively
- Can reach critical point much faster than gradient descent
 - But useful only when nearby point is a minimum

Summary of Gradient Methods

- First order optimization algorithms: those that use only the gradient
- Second order optimization algorithms: use the Hessian matrix such as Newton's method
- Family of functions used in ML is complicated, so optimization is more complex than in other fields
 - No guarantees
- Some guarantees by using Lipschitz continuous functions,

$$|f(\mathbf{x}) - f(\mathbf{y})| \leq L \|\mathbf{x} - \mathbf{y}\|_2$$

- with Lipschitz constant L

Convex Optimization

- Applicable only to convex functions – functions which are well-behaved,
 - e.g., lack saddle points and all local minima are global minima
- For such functions, Hessian is positive semi-definite everywhere
- Many ML optimization problems, particularly deep learning, cannot be expressed as convex optimization

Constrained Optimization

- We may wish to optimize $f(\mathbf{x})$ when the solution \mathbf{x} is constrained to lie in set S
 - Such values of \mathbf{x} are feasible solutions
- Often we want a solution that is small, such as $||\mathbf{x}|| \leq 1$
- Simple approach: modify gradient descent taking constraint into account (using Lagrangian formulation)

Ex: Least squares with Lagrangian

- We wish to minimize $f(\mathbf{x}) = \frac{1}{2} \|A\mathbf{x} - \mathbf{b}\|^2$
 - Subject to constraint $\mathbf{x}^T \mathbf{x} \leq 1$
- We introduce the Lagrangian $L(\mathbf{x}, \lambda) = f(\mathbf{x}) + \lambda(\mathbf{x}^T \mathbf{x} - 1)$
 - And solve the problem $\min_{\mathbf{x}} \max_{\lambda, \lambda \geq 0} L(\mathbf{x}, \lambda)$
- For the unconstrained problem (no Lagrangian) the smallest norm solution is $\mathbf{x} = A^+ \mathbf{b}$
 - If this solution is not feasible, differentiate Lagrangian wrt \mathbf{x} to obtain $A^T A \mathbf{x} - A^T \mathbf{b} + 2\lambda \mathbf{x} = 0$
 - Solution takes the form $\mathbf{x} = (A^T A + 2\lambda I)^{-1} A^T \mathbf{b}$
 - Choosing λ : continue solving linear equation and increasing λ until \mathbf{x} has the correct norm

Generalized Lagrangian: KKT

- More sophisticated than Lagrangian
- Karush-Kuhn-Tucker is a very general solution to constrained optimization
- While Lagrangian allows equality constraints, KKT allows both equality and inequality constraints
- To define a generalized Lagrangian we need to describe S in terms of equalities and inequalities

Generalized Lagrangian

- Set S is described in terms of m functions $g(i)$ and n functions $h(j)$ so that

$$S = \left\{ \mathbf{x} \mid \forall i, g^{(i)}(\mathbf{x}) = 0 \text{ and } \forall j, h^{(j)}(\mathbf{x}) \leq 0 \right\}$$

- Functions of g are equality constraints and functions of h are inequality constraints
- Introduce new variables λ_i and α_j for each constraint (called KKT multipliers) giving the generalized Lagrangian

$$L(\mathbf{x}, \lambda, \alpha) = f(\mathbf{x}) + \sum_i \lambda_i g^{(i)}(\mathbf{x}) + \sum_j \alpha_j h^{(j)}(\mathbf{x})$$

- We can now solve the unconstrained optimization problem